

Cyber Security Mitigation and Response AI

Dataset Source: <https://www.kaggle.com/sampadab17/network-intrusion-detection>
(<https://www.kaggle.com/sampadab17/network-intrusion-detection>).

In [75]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sweetviz as sv
from scipy import stats
import seaborn as sns
```

In [76]:

```
test_cs = pd.read_csv('test_data.csv')
train_cs = pd.read_csv('train_data.csv')
```


In [77]:

```
train_cs
```

Out[77]:

| | duration | protocol_type | service | flag | src_bytes | dst_bytes | land | wrong_fragment | ui |
|-------|----------|---------------|----------|------|-----------|-----------|------|----------------|----|
| 0 | 0 | tcp | ftp_data | SF | 491 | 0 | 0 | 0 | |
| 1 | 0 | udp | other | SF | 146 | 0 | 0 | 0 | |
| 2 | 0 | tcp | private | S0 | 0 | 0 | 0 | 0 | |
| 3 | 0 | tcp | http | SF | 232 | 8153 | 0 | 0 | |
| 4 | 0 | tcp | http | SF | 199 | 420 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 25187 | 0 | tcp | exec | RSTO | 0 | 0 | 0 | 0 | |
| 25188 | 0 | tcp | ftp_data | SF | 334 | 0 | 0 | 0 | |
| 25189 | 0 | tcp | private | REJ | 0 | 0 | 0 | 0 | |
| 25190 | 0 | tcp | nnsp | S0 | 0 | 0 | 0 | 0 | |
| 25191 | 0 | tcp | finger | S0 | 0 | 0 | 0 | 0 | |

25192 rows × 42 columns




In [78]:

```
train_cs.describe()
```

Out[78]:

| | duration | src_bytes | dst_bytes | land | wrong_fragment | urgent |
|-------|--------------|--------------|--------------|--------------|----------------|--------------|
| count | 25192.000000 | 2.519200e+04 | 2.519200e+04 | 25192.000000 | 25192.000000 | 25192.000000 |
| mean | 305.054104 | 2.433063e+04 | 3.491847e+03 | 0.000079 | 0.023738 | 0.00004 |
| std | 2686.555640 | 2.410805e+06 | 8.883072e+04 | 0.008910 | 0.260221 | 0.00630 |
| min | 0.000000 | 0.000000e+00 | 0.000000e+00 | 0.000000 | 0.000000 | 0.00000 |
| 25% | 0.000000 | 0.000000e+00 | 0.000000e+00 | 0.000000 | 0.000000 | 0.00000 |
| 50% | 0.000000 | 4.400000e+01 | 0.000000e+00 | 0.000000 | 0.000000 | 0.00000 |
| 75% | 0.000000 | 2.790000e+02 | 5.302500e+02 | 0.000000 | 0.000000 | 0.00000 |
| max | 42862.000000 | 3.817091e+08 | 5.151385e+06 | 1.000000 | 3.000000 | 1.00000 |

8 rows × 38 columns



In [79]:

```
train_cs.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25192 entries, 0 to 25191
Data columns (total 42 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   duration                             25192 non-null  int64
1   protocol_type                        25192 non-null  object
2   service                              25192 non-null  object
3   flag                                 25192 non-null  object
4   src_bytes                            25192 non-null  int64
5   dst_bytes                            25192 non-null  int64
6   land                                 25192 non-null  int64
7   wrong_fragment                       25192 non-null  int64
8   urgent                               25192 non-null  int64
9   hot                                  25192 non-null  int64
10  num_failed_logins                    25192 non-null  int64
11  logged_in                            25192 non-null  int64
12  num_compromised                      25192 non-null  int64
13  root_shell                           25192 non-null  int64
14  su_attempted                         25192 non-null  int64
15  num_root                             25192 non-null  int64
16  num_file_creations                  25192 non-null  int64
17  num_shells                           25192 non-null  int64
18  num_access_files                    25192 non-null  int64
19  num_outbound_cmds                   25192 non-null  int64
20  is_host_login                       25192 non-null  int64
21  is_guest_login                       25192 non-null  int64
22  count                               25192 non-null  int64
23  srv_count                           25192 non-null  int64
24  serror_rate                         25192 non-null  float64
25  srv_serror_rate                     25192 non-null  float64
26  rerror_rate                         25192 non-null  float64
27  srv_rerror_rate                     25192 non-null  float64
28  same_srv_rate                       25192 non-null  float64
29  diff_srv_rate                       25192 non-null  float64
30  srv_diff_host_rate                  25192 non-null  float64
31  dst_host_count                       25192 non-null  int64
32  dst_host_srv_count                  25192 non-null  int64
33  dst_host_same_srv_rate              25192 non-null  float64
34  dst_host_diff_srv_rate              25192 non-null  float64
35  dst_host_same_src_port_rate         25192 non-null  float64
36  dst_host_srv_diff_host_rate         25192 non-null  float64
37  dst_host_serror_rate                25192 non-null  float64
38  dst_host_srv_serror_rate            25192 non-null  float64
39  dst_host_rerror_rate                25192 non-null  float64
40  dst_host_srv_rerror_rate            25192 non-null  float64
41  class                               25192 non-null  object
dtypes: float64(15), int64(23), object(4)
memory usage: 8.1+ MB
```

In [80]:

```
#demotrain = sv.analyze(train_cs)
#demotrain.show_html()
```

In [81]:

```
train_cs.isnull().sum()
```

Out[81]:

```
duration          0
protocol_type     0
service           0
flag              0
src_bytes         0
dst_bytes         0
land              0
wrong_fragment    0
urgent            0
hot               0
num_failed_logins 0
logged_in         0
num_compromised   0
root_shell        0
su_attempted      0
num_root          0
num_file_creations 0
num_shells         0
num_access_files  0
num_outbound_cmds 0
is_host_login     0
is_guest_login    0
count             0
srv_count         0
serror_rate       0
srv_serror_rate   0
rerror_rate       0
srv_rerror_rate   0
same_srv_rate     0
diff_srv_rate     0
srv_diff_host_rate 0
dst_host_count    0
dst_host_srv_count 0
dst_host_same_srv_rate 0
dst_host_diff_srv_rate 0
dst_host_same_src_port_rate 0
dst_host_srv_diff_host_rate 0
dst_host_serror_rate 0
dst_host_srv_serror_rate 0
dst_host_rerror_rate 0
dst_host_srv_rerror_rate 0
class             0
dtype: int64
```

In [82]:

```
train_cs.duplicated().sum()
```

Out[82]:


0

In [83]:

```
sns.countplot(data=train_cs, x='class')
```

Out[83]:

<AxesSubplot:xlabel='class', ylabel='count'>



In [84]:

```
(train_cs['class'].values == 'anomaly').sum()
```

Out[84]:


11743

In [85]:

```
sns.countplot(data=train_cs, x='protocol_type')
```

Out[85]:

<AxesSubplot:xlabel='protocol_type', ylabel='count'>



In [86]:

```
train_cs['service'].unique()
```

Out[86]:


```
array(['ftp_data', 'other', 'private', 'http', 'remote_job', 'name',
      'netbios_ns', 'eco_i', 'mtp', 'telnet', 'finger', 'domain_u',
      'supdup', 'uucp_path', 'Z39_50', 'smtp', 'csnet_ns', 'uucp',
      'netbios_dgm', 'urp_i', 'auth', 'domain', 'ftp', 'bgp', 'ldap',
      'ecr_i', 'gopher', 'vmnet', 'sysstat', 'http_443', 'efs', 'whois',
      'imap4', 'iso_tsap', 'echo', 'klogin', 'link', 'sunrpc', 'login',
      'kshell', 'sql_net', 'time', 'hostnames', 'exec', 'ntp_u',
      'discard', 'nntp', 'courier', 'ctf', 'ssh', 'daytime', 'shell',
      'netstat', 'pop_3', 'nnsf', 'IRC', 'pop_2', 'printer', 'tim_i',
      'pm_dump', 'red_i', 'netbios_ssn', 'rje', 'X11', 'urh_i',
      'http_8001'], dtype=object)
```

In [87]:

```
sns.countplot(data=train_cs, x='flag')
```

Out[87]:

```
<AxesSubplot:xlabel='flag', ylabel='count'>
```



In [88]:

```
nor_data = train_cs.loc[train_cs['class'] == 'normal']
```


In [89]:

```
nor_data #normal data distribution
```

Out[89]:

| | duration | protocol_type | service | flag | src_bytes | dst_bytes | land | wrong_fragment | urg |
|-------|----------|---------------|----------|------|-----------|-----------|------|----------------|-----|
| 0 | 0 | tcp | ftp_data | SF | 491 | 0 | 0 | 0 | |
| 1 | 0 | udp | other | SF | 146 | 0 | 0 | 0 | |
| 3 | 0 | tcp | http | SF | 232 | 8153 | 0 | 0 | |
| 4 | 0 | tcp | http | SF | 199 | 420 | 0 | 0 | |
| 12 | 0 | tcp | http | SF | 287 | 2251 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 25176 | 0 | tcp | ftp_data | SF | 748 | 0 | 0 | 0 | |
| 25177 | 0 | tcp | http | SF | 293 | 2486 | 0 | 0 | |
| 25184 | 29 | tcp | ftp | SF | 329 | 1063 | 0 | 0 | |
| 25185 | 1 | tcp | smtp | SF | 2896 | 333 | 0 | 0 | |
| 25186 | 0 | tcp | http | S1 | 339 | 14600 | 0 | 0 | |

13449 rows × 42 columns



In [90]:

```
ano_data = train_cs.loc[train_cs['class'] == 'anomaly']
```

In [91]:

```
ano_data #anomaly data distribution
```

Out[91]:

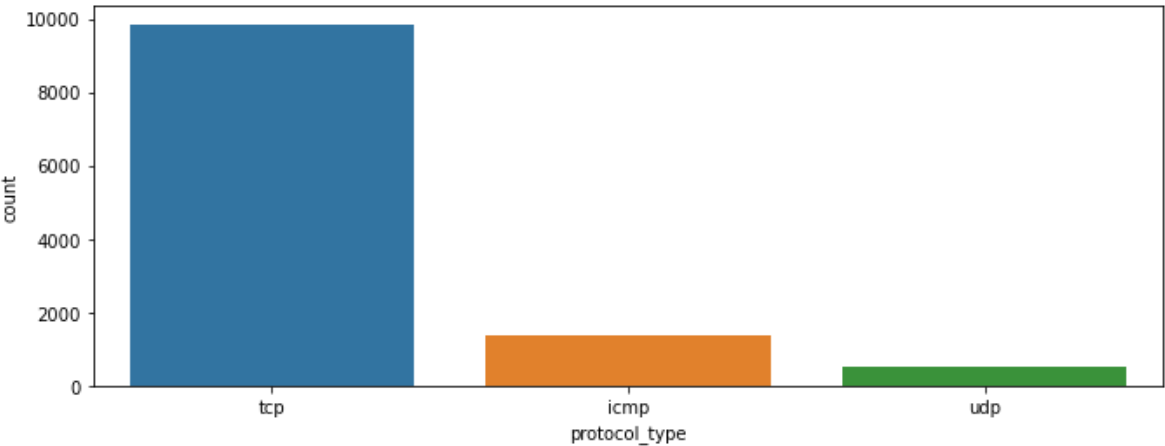
| e | dst_host_same_src_port_rate | dst_host_srv_diff_host_rate | dst_host_serror_rate | dst_host_srv_sei |
|----|-----------------------------|-----------------------------|----------------------|------------------|
| 5 | 0.00 | 0.00 | 1.0 | |
| 7 | 0.00 | 0.00 | 0.0 | |
| 5 | 0.00 | 0.00 | 1.0 | |
| 7 | 0.00 | 0.00 | 1.0 | |
| 5 | 0.00 | 0.00 | 1.0 | |
| .. | ... | ... | ... | |
| 6 | 0.00 | 0.00 | 0.0 | |
| 0 | 1.00 | 0.18 | 0.0 | |
| 7 | 0.00 | 0.00 | 0.0 | |
| 6 | 0.00 | 0.00 | 1.0 | |
| 3 | 0.01 | 0.00 | 1.0 | |

In [92]:

```
sns.countplot(data=ano_data, x='protocol_type')
```

Out[92]:

<AxesSubplot:xlabel='protocol_type', ylabel='count'>




In anomaly data tcp and icmp protocols have higher usages

In [93]:

```
sns.countplot(data=nor_data, x='protocol_type')
```

Out[93]:

<AxesSubplot:xlabel='protocol_type', ylabel='count'>




In normal distribution tcp and udp protocols have higher usage

In [94]:

```
sns.countplot(data=ano_data, x='flag')
```

Out[94]:

<AxesSubplot:xlabel='flag', ylabel='count'>




In [95]:

```
sns.countplot(data=nor_data, x='flag')
```

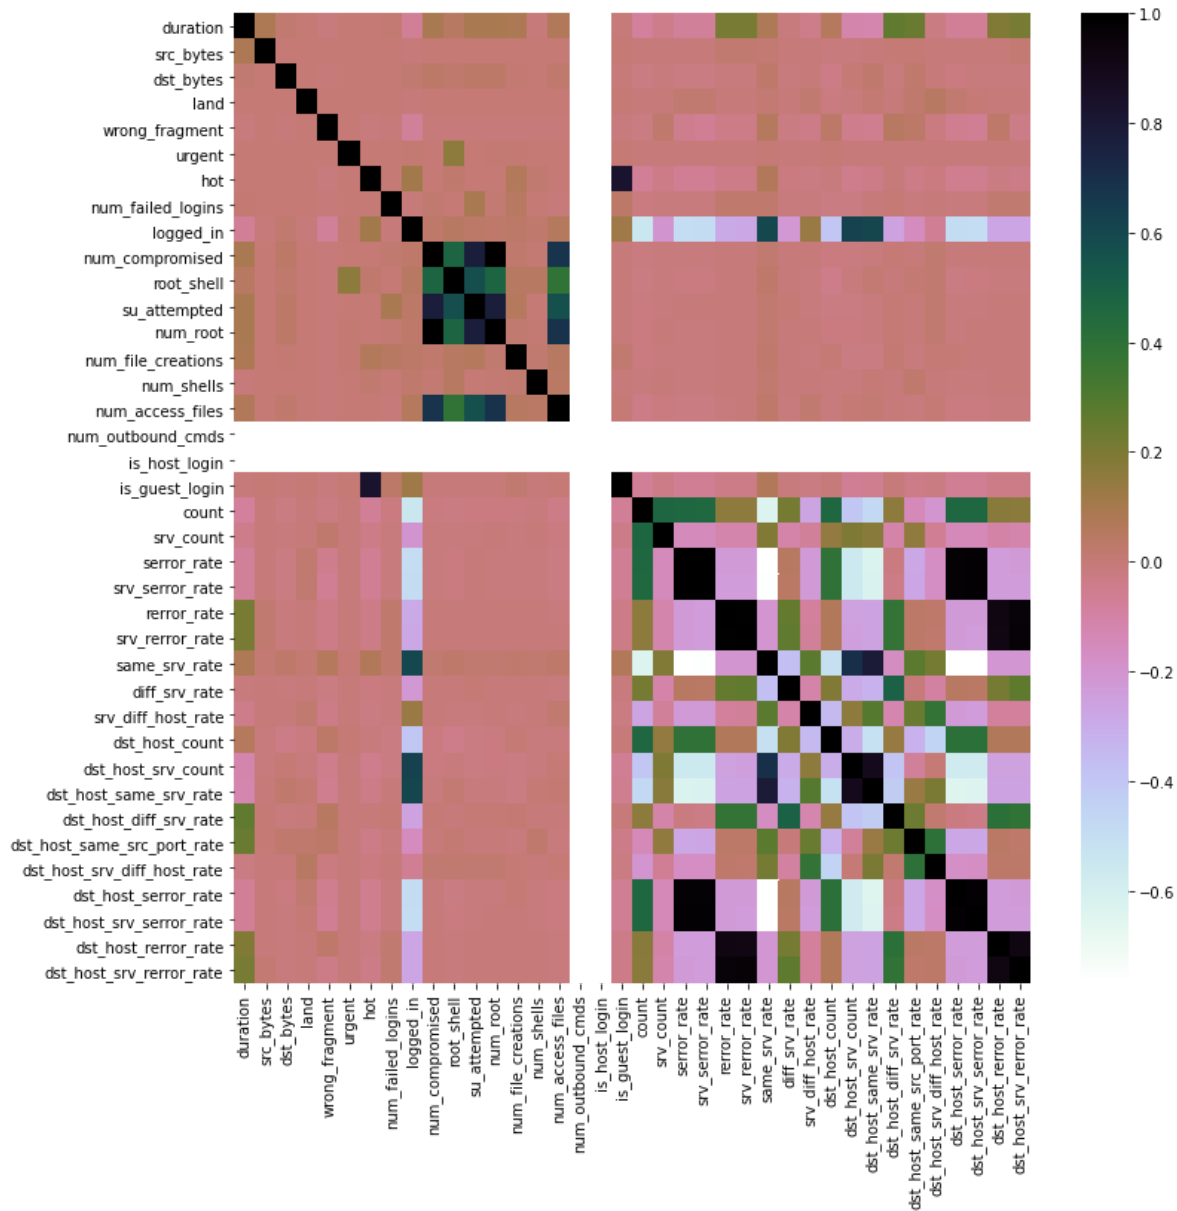
Out[95]:

<AxesSubplot:xlabel='flag', ylabel='count'>



In [96]:

```
corr = train_cs.corr()
plt.figure(figsize=(12,12))
sns.heatmap(corr, cmap='cubehelix_r');
```



In [97]:

```
train_cs.drop(['num_outbound_cmds', 'is_host_login'], axis=1, inplace=True)
test_cs.drop(['num_outbound_cmds', 'is_host_login'], axis=1, inplace=True)
```

Separating numerical and object columns

Combining all Numerical Columns

In [98]:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

cols = train_cs.select_dtypes(include=['float64', 'int64']).columns
sc_train = scaler.fit_transform(train_cs.select_dtypes(include=['float64', 'int64']))
sc_test = scaler.fit_transform(test_cs.select_dtypes(include=['float64', 'int64']))

train_num = pd.DataFrame(sc_train, columns = cols)
test_num = pd.DataFrame(sc_test, columns = cols)
```

Finding Object type Columns

In [99]:

```
from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()

# extract object attributes from both train and test sets
train_obj = train_cs.select_dtypes(include=['object']).copy()
test_obj = test_cs.select_dtypes(include=['object']).copy()

# encode the object attributes
train_a = train_obj.apply(encoder.fit_transform)
test_a = test_obj.apply(encoder.fit_transform)

# separate 'class' column from encoded data
train_noclass = train_a.drop(['class'], axis=1)
Ytrain_class = train_a[['class']].copy()
```

In [100]:

```
train_x = pd.concat([train_num,train_noclass],axis=1)
train_y = train_cs['class']
train_x.shape
```

Out[100]:

(25192, 39)

In [101]:

```
test_ = pd.concat([test_num,test_a],axis=1)
test_.shape
```

Out[101]:

(22544, 39)

Feature Selection

In [102]:

```


from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier();

# fit random forest classifier on the training set
rfc.fit(train_x, train_y);

# extract important features
score = np.round(rfc.feature_importances_,3)
imp = pd.DataFrame({'feature':train_x.columns,'importance':score})
imp = imp.sort_values('importance',ascending=False).set_index('feature')

# plot importances
plt.rcParams['figure.figsize'] = (11, 4)
imp.plot.bar();

```



In [103]:

```
from sklearn.feature_selection import RFE
import itertools
rfc = RandomForestClassifier()

# create the RFE model and select 15 attributes
rfe = RFE(rfc, n_features_to_select=15)
rfe = rfe.fit(train_x, train_y)

# summarize the selection of the attributes
feature_map = [(i, v) for i, v in itertools.zip_longest(rfe.get_support(), train_x.columns)]
features = [v for i, v in feature_map if i==True]

features
```

Out[103]:

```
['src_bytes',
 'dst_bytes',
 'logged_in',
 'count',
 'srv_count',
 'same_srv_rate',
 'diff_srv_rate',
 'dst_host_srv_count',
 'dst_host_same_srv_rate',
 'dst_host_diff_srv_rate',
 'dst_host_same_src_port_rate',
 'dst_host_srv_diff_host_rate',
 'protocol_type',
 'service',
 'flag']
```

In [104]:

```
train_feature = train_x[features]
train_feature
```

Out[104]:

| src_rate | dst_host_srv_count | dst_host_same_srv_rate | dst_host_diff_srv_rate | dst_host_same_src_port_rate |
|----------|--------------------|------------------------|------------------------|-----------------------------|
| 349282 | -0.813985 | -0.779157 | -0.280673 | 0.0 |
| 490836 | -1.030895 | -1.157831 | 2.764403 | 2.0 |
| 342773 | -0.804947 | -0.935081 | -0.173828 | -0.0 |
| 349282 | 1.264742 | 1.069663 | -0.440940 | -0.0 |
| 349282 | 1.264742 | 1.069663 | -0.440940 | -0.0 |
| ... | ... | ... | ... | ... |
| 342773 | -0.976667 | -1.091006 | -0.120406 | -0.0 |
| 349282 | -0.687453 | 1.069663 | -0.440940 | 2.0 |
| 342773 | -0.922440 | -1.046456 | -0.066984 | -0.0 |
| 313235 | -0.859174 | -0.979631 | -0.120406 | -0.0 |
| 266804 | -0.597074 | -0.734607 | -0.280673 | -0.0 |

Defining Train and Test

In [105]:

```
from sklearn.model_selection import train_test_split

X_train, X_test, Y_train, Y_test = train_test_split(train_feature, train_y, test_size = 0.2)
```

Applying Logistic Regression

In [106]:

```
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.model_selection import cross_val_score
```

In [107]:

```
LGR_Classifier = LogisticRegression(n_jobs=-1, random_state=0)
LGR_Classifier.fit(X_train, Y_train);
```


In [122]:

```
LGR_Classifier.coef_
```

Out[122]:

```
array([[ 9.81364175e-02, -3.65668863e-01, -7.22050813e-02,  
        -1.65244595e-03, -1.31498884e+00, -9.28785294e-02,  
        -5.90453266e-01, -2.05501508e-02,  1.91271626e-01,  
        -4.50808136e-01, -7.36916857e-02,  3.45145878e-01,  
         3.01212944e-01,  3.97918465e-02,  1.18981959e-01,  
         2.18297502e-02,  3.97770604e-01, -1.48222618e+00,  
         3.98018259e-01, -4.21418384e-01, -1.31829496e+00,  
        -4.56260289e-01, -1.13789835e+00,  1.19844649e+00,  
         3.55295807e-01, -2.63479517e-01, -9.60747615e-01,  
         1.70019090e+00, -1.20766337e+00, -3.01584356e-01,  
        -8.48080088e-01, -3.32264621e-01, -5.11430962e-01,  
        -9.89170359e-01, -3.09202230e-01,  1.07031839e-01,  
         1.60210425e+00,  4.85916741e-03, -3.65035186e-01]])
```

Evaluating Model

In [108]:

```

models = []
models.append(('LogisticRegression', LGR_Classifier))

for i, v in models:
    scores = cross_val_score(v, X_train, Y_train, cv=10)
    accuracy = metrics.accuracy_score(Y_train, v.predict(X_train))
    confusion_matrix = metrics.confusion_matrix(Y_train, v.predict(X_train))
    classification = metrics.classification_report(Y_train, v.predict(X_train))
    print()
    print('===== {} Model Evaluation =====')
    print()
    print("Cross Validation Mean Score:" "\n", scores.mean())
    print()
    print("Model Accuracy:" "\n", accuracy)
    print()
    print("Confusion matrix:" "\n", confusion_matrix)
    print()
    print("Classification report:" "\n", classification)
    print()

```

```

===== LogisticRegression Model Evaluation =====
=====

```

```

Cross Validation Mean Score:
0.943632391371936

```

```

Model Accuracy:
0.9439504604636393

```

```

Confusion matrix:
[[8162  652]
 [ 407 9673]]

```

```

Classification report:

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| anomaly | 0.95 | 0.93 | 0.94 | 8814 |
| normal | 0.94 | 0.96 | 0.95 | 10080 |
| accuracy | | | 0.94 | 18894 |
| macro avg | 0.94 | 0.94 | 0.94 | 18894 |
| weighted avg | 0.94 | 0.94 | 0.94 | 18894 |

Validating Model

In [109]:

```

for i, v in models:
    accuracy = metrics.accuracy_score(Y_test, v.predict(X_test))
    confusion_matrix = metrics.confusion_matrix(Y_test, v.predict(X_test))
    classification = metrics.classification_report(Y_test, v.predict(X_test))
    print()
    print('===== {} Model Test Results =====')
    print()
    print("Model Accuracy:" "\n", accuracy)
    print()
    print("Confusion matrix:" "\n", confusion_matrix)
    print()
    print("Classification report:" "\n", classification)
    print()

```

```

===== LogisticRegression Model Test Results =====
=====

```

```

Model Accuracy:
0.9402985074626866

```

```

Confusion matrix:
[[2703  226]
 [ 150 3219]]

```

```

Classification report:

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| anomaly | 0.95 | 0.92 | 0.93 | 2929 |
| normal | 0.93 | 0.96 | 0.94 | 3369 |
| accuracy | | | 0.94 | 6298 |
| macro avg | 0.94 | 0.94 | 0.94 | 6298 |
| weighted avg | 0.94 | 0.94 | 0.94 | 6298 |

Another approach to improve model's accuracy rate

In [110]:

```

from sklearn.model_selection import train_test_split

X_train, X_test, Y_train, Y_test = train_test_split(train_x, train_y, test_size = 0.25, ran

```

In [111]:

```

LGR_Classifier = LogisticRegression(n_jobs=-1, random_state=0)
LGR_Classifier.fit(X_train, Y_train);

```

Evaluation Model

In [112]:

```

models = []
models.append(('LogisticRegression', LGR_Classifier))

for i, v in models:
    scores = cross_val_score(v, X_train, Y_train, cv=10)
    accuracy = metrics.accuracy_score(Y_train, v.predict(X_train))
    confusion_matrix = metrics.confusion_matrix(Y_train, v.predict(X_train))
    classification = metrics.classification_report(Y_train, v.predict(X_train))
    print()
    print('===== {} Model Evaluation =====')
    print()
    print("Cross Validation Mean Score:" "\n", scores.mean())
    print()
    print("Model Accuracy:" "\n", accuracy)
    print()
    print("Confusion matrix:" "\n", confusion_matrix)
    print()
    print("Classification report:" "\n", classification)
    print()

```

```

===== LogisticRegression Model Evaluation =====
=====

```

```

Cross Validation Mean Score:
0.9547470036776546

```

```

Model Accuracy:
0.9556472954377051

```

```

Confusion matrix:
[[8298  516]
 [ 322 9758]]

```

```

Classification report:

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| anomaly | 0.96 | 0.94 | 0.95 | 8814 |
| normal | 0.95 | 0.97 | 0.96 | 10080 |
| accuracy | | | 0.96 | 18894 |
| macro avg | 0.96 | 0.95 | 0.96 | 18894 |
| weighted avg | 0.96 | 0.96 | 0.96 | 18894 |

Validation Model

In [113]:

```

for i, v in models:
    accuracy = metrics.accuracy_score(Y_test, v.predict(X_test))
    confusion_matrix = metrics.confusion_matrix(Y_test, v.predict(X_test))
    classification = metrics.classification_report(Y_test, v.predict(X_test))
    print()
    print('===== {} Model Test Results =====')
    print()
    print("Model Accuracy:" "\n", accuracy)
    print()
    print("Confusion matrix:" "\n", confusion_matrix)
    print()
    print("Classification report:" "\n", classification)
    print()

```

```

===== LogisticRegression Model Test Results =====
=====

```

```

Model Accuracy:
0.9537948555096856

```

```

Confusion matrix:
[[2754 175]
 [ 116 3253]]

```

```

Classification report:

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| anomaly | 0.96 | 0.94 | 0.95 | 2929 |
| normal | 0.95 | 0.97 | 0.96 | 3369 |
| accuracy | | | 0.95 | 6298 |
| macro avg | 0.95 | 0.95 | 0.95 | 6298 |
| weighted avg | 0.95 | 0.95 | 0.95 | 6298 |

Predicting normal and anomaly behaviour on Test dataset

In [114]:

```
pred_log = LGR_Classifier.predict(test_)
```

In [115]:

```
pred_log
```

Out[115]:

```

array(['anomaly', 'anomaly', 'normal', ..., 'normal', 'normal', 'anomaly'],
      dtype=object)

```


In [116]:

test_cs

Out[116]:

| | duration | protocol_type | service | flag | src_bytes | dst_bytes | land | wrong_fragment |
|-------|----------|---------------|----------|------|-----------|-----------|------|----------------|
| 0 | 0 | tcp | private | REJ | 0 | 0 | 0 | 0 |
| 1 | 0 | tcp | private | REJ | 0 | 0 | 0 | 0 |
| 2 | 2 | tcp | ftp_data | SF | 12983 | 0 | 0 | 0 |
| 3 | 0 | icmp | eco_i | SF | 20 | 0 | 0 | 0 |
| 4 | 1 | tcp | telnet | RSTO | 0 | 15 | 0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 22539 | 0 | tcp | smtp | SF | 794 | 333 | 0 | 0 |
| 22540 | 0 | tcp | http | SF | 317 | 938 | 0 | 0 |
| 22541 | 0 | tcp | http | SF | 54540 | 8314 | 0 | 0 |
| 22542 | 0 | udp | domain_u | SF | 42 | 42 | 0 | 0 |
| 22543 | 0 | tcp | sunrpc | REJ | 0 | 0 | 0 | 0 |

22544 rows × 39 columns



In [117]:

```
pred_df = pd.DataFrame(pred_log, columns = ['class'])
test_output = pd.concat([test_cs, pred_df],axis=1)
```

In [118]:

```
test_output
```

Out[118]:


| e | dst_host_same_src_port_rate | dst_host_srv_diff_host_rate | dst_host_serror_rate | dst_host_srv_sei |
|----|-----------------------------|-----------------------------|----------------------|------------------|
| 6 | 0.00 | 0.00 | 0.00 | |
| 6 | 0.00 | 0.00 | 0.00 | |
| 4 | 0.61 | 0.02 | 0.00 | |
| 0 | 1.00 | 0.28 | 0.00 | |
| 7 | 0.03 | 0.02 | 0.00 | |
| .. | ... | ... | ... | |
| 6 | 0.01 | 0.01 | 0.01 | |
| 0 | 0.01 | 0.01 | 0.01 | |
| 0 | 0.00 | 0.00 | 0.00 | |
| 1 | 0.00 | 0.00 | 0.00 | |
| 3 | 0.00 | 0.00 | 0.00 | |

In [119]:

```
sns.countplot(data = test_output, x = 'class')
```

Out[119]:

<AxesSubplot:xlabel='class', ylabel='count'>



In []:

